

VeriSigil AI

PRODUCT BLUEPRINT

The Constitutional Execution Layer for Autonomous AI

Architecture Freeze · Engineering Direction · 90-Day Roadmap

*"Every AI action deserves
authorization before execution."*

VeriSigil AI

raheem@verisigil.ai · verisigil.ai

Version 1.0 · August 2026 · CONFIDENTIAL

CORE MISSION

Before any autonomous AI action executes, VeriSigil AI verifies that the AI has the authority to perform that action.

Every feature must strengthen this mission. Everything else is out of scope.

The Golden Engineering Rule

Every proposed feature must pass this filter:

Does this feature help determine whether an autonomous AI action should legitimately execute **before execution**?

YES → Build it

NO → Reject it or move it into a future independent product

Product Identity

WE ARE

- Execution Authorization Infrastructure
- The layer between AI intent and real-world consequence
- Think: Stripe for AI Execution Trust
- Think: Okta for AI Identity & Authority
- Think: Cloudflare for AI Edge Governance

WE ARE NOT

- AI Chat / Copilot / Assistant
- CRM or Business Automation Platform
- AI Consulting Company / CEO Digital Twin
- Workflow Builder / Sales Automation
- ERP / Accounting / Project Management

ONE SENTENCE — USE EVERYWHERE

"Every AI action deserves authorization before execution."

Architecture: 5 Meta-Layers

LAYER 1 Organizational Intelligence

Purpose: Understand what exists before governing it.

BUILD

- Organization Graph
- Human Role Discovery
- AI Asset Discovery
- Process Discovery
- Decision Graph
- Knowledge Graph

NOT

- Revenue audits
- CRM analytics
- General consulting

LAYER 2 Delegation Intelligence

Purpose: Design authority. Not replace authority.

BUILD

- Decision Routing (Jethro Hierarchy)
- Delegation Passport
- Capable People Registry
- Trust Capacity Score
- Delegation Health

NOT

- HR management
- Payroll software
- Business intelligence

LAYER 3 Constitutional Governance

Purpose: Decide whether execution should occur.

BUILD

- Runtime Gateway (6-Gate)
- Policy Engine
- Evidence Engine
- Challenge & Appeal
- Drift Detection
- Audit Ledger

NOT

- Auto policy deployment
- Auto appeal resolution
- Auto amendments

LAYER 4 Autonomous Operations

Purpose: Govern AI workers. Not build AI workers.

BUILD

- AI Workforce Registry
- Voice Integrity Layer
- Runtime Monitoring
- Trust Mesh / IATP
- Credential Governance
- Boundary Register

NOT

- Voice AI company
- Sales AI
- Customer support platform

LAYER 5 Organizational Evolution

Purpose: Organizations change. Governance evolves.

BUILD

- Constitutional Memory
- Policy Gap Detection
- Authority Simulation
- Governance DNA
- Policy Evolution

NOT

- Strategy consulting
- Market research
- Financial forecasting

Engineering Priority: Next 90 Days

PHASE 1 — Immediate

- VeriSigil Runtime API — intercept, policy, human approval, evidence
- Basic Dashboard — per-agent governance health
- Demo Application — guided journey + governance receipt
- Python SDK on PyPI — five lines to integrate

PHASE 2 — This Month

- Delegation Passport — allowed/forbidden/expiry/budget/supervisor
- Decision Routing Engine — Jethro hierarchy tier routing
- Voice Integrity Layer — Bland, Retell, Vapi, ElevenLabs
- Trust Mesh / IATP — inter-agent trust negotiation
- Boundary Register + Credential Governance

PHASE 3 — Next Quarter

- Constitutional Memory — permanent append-only decision record
- Governance DNA — organizational risk philosophy
- JavaScript SDK · Java SDK · Enterprise Connectors (LangChain)
- Independent Security Assessment / Penetration Test

What VeriSigil Will Never Build

OUT OF SCOPE — NO EXCEPTIONS

CEO Digital Twin · CRM · Sales Automation · AI Chat / Assistant
General Workflow Automation · Accounting · ERP · Marketing Platform
Business Consulting · Auto policy deployment · Auto appeal resolution

Current Platform Status

Production: verisigil-api-production.up.railway.app · ~933 endpoints · 76,819 lines

Item	Status	Evidence
6-Gate Runtime Intercept	LIVE	POST /v1/intercept
Ed25519 Receipt Signing	LIVE	CLARA 25/25 offline verified
Authorization Object (AO)	LIVE	Run 3b: 1 PROCEED + 4 HALT concurrent
SQLite Replay Prevention	LIVE	ALREADY_CONSUMED across 5h + restarts
Delegation Passport	LIVE	POST /v1/delegation/passport
Decision Routing Engine	LIVE	POST /v1/decision/route (Jethro 6-tier)
Voice Integrity Layer	LIVE	POST /v1/voice/session/start

IATP Inter-Agent Trust	LIVE	POST /v1/trust/negotiate
AI Immune System	LIVE	POST /v1/immune/observe + quarantine
VGES Open Benchmark	LIVE	GET /v1/benchmark/vges (CC BY 4.0)
OMNIX IWA	IN REVIEW	EWP-IWA-VERISIGIL-001 v1.1
Independent Validation	RUN 3 CLOSED	Alkama Eqbal 2026-08-07
Verification Report #001	PUBLISHED	verisigilai.com/verification-reports.html
Python SDK	READY	verisigil_ai_sdk.py — pending PyPI

LONG-TERM SUCCESS METRIC

Enterprise architects automatically think of VeriSigil AI whenever they ask:

"Before this AI acts, who verifies that it has the authority to act?"

Identity = Okta · Payments = Stripe · AI Execution Trust = VeriSigil